

UNIVERSITY OF SOUTHERN DENMARK

DEPARTMENT OF MATHEMATICS AND COMPUTER SCIENCE

SPD801: MASTER THESIS IN COMPUTER SCIENCE

---

# Formalising $\pi$ LL in Lean

---

*Author*

Mikkel Brix Nielsen  
mikke21

*Supervisors*

Supervisor: Fabrizio Montesi  
Cosupervisor: Marco Peressotti  
Cosupervisor: Xueying Qin

18th May 2026



## 1 Introduction

The Curry-Howard correspondence establishes a deep connection between logic and computation, where proofs correspond to programs and logical propositions correspond to types. This perspective has been particularly influential in the study of concurrent systems, where linear logic and session types provide a logical foundation for structured communication.

In this setting, linear logic captures resource-sensitive reasoning, while session types describe communication protocols between concurrent processes. The resulting correspondence, often referred to as propositions-as-sessions, provides a principled interpretation of interaction as logical proof transformation, and has been further related to process calculi such as the  $\pi$ -calculus.

This line of work has been further developed by [Montesi and Peressotti 2021], who provide a unified account reconciling linear logic, the  $\pi$ -calculus, and their metatheory. Their framework serves as the formal basis for the system  $\pi$ LL considered in this thesis.

*Motivation and correctness concerns.* Despite the elegance of these theoretical correspondences, their informal and paper-based presentations are prone to subtle errors. In particular, the application of inference rules in linear logic and session-typed systems often relies on implicit ordering conventions and intuitive interpretations of syntax. This can lead to mistakes that are not immediately apparent in written derivations.

As an illustrative example, small inconsistencies can arise in the presentation of typing rules within such systems, where the intuitive alignment between variable ordering and type assignment does not precisely match the formal rule structure. These discrepancies are not flaws of the underlying theory, but rather reflect the difficulty of maintaining precision in informal presentations, including in the framework of Montesi and Peressotti.

As observed in discussions with the authors of the framework, such issues are easy to overlook because rule application is often guided by intuition rather than strict syntactic structure. This highlights that even formally correct rules may be unintuitive to apply due to their representation.

These observations motivate a more disciplined approach to formal reasoning, where correctness is ensured by construction rather than by inspection.

*Why mechanization.* Mechanization provides a means of bridging this gap. Rather than extending the underlying theory, the goal of mechanization in this work is to validate and internalize existing results within a proof assistant. This allows us to eliminate ambiguity arising from informal presentation and to ensure that all derivations are checked with respect to a precise formal semantics.

Moreover, mechanized developments have repeatedly demonstrated their ability to uncover small but meaningful inconsistencies in paper proofs, supporting the view that proof assistants serve as a valuable tool for verifying, rather than merely implementing, theoretical frameworks.

*Contributions.* In this thesis, we present a mechanized formalization of the  $\pi$ -calculus-based linear logic system  $\pi$ LL as developed by Montesi and Peressotti, with particular emphasis on its multiplicative fragment. The contributions are as follows:

- A formalization of the syntax and typing system of  $\pi$ LL in Lean, including processes, session types, and substitution operations.
- A formalization of the operational semantics for the multiplicative fragment of  $\pi$ LL.
- Formal proofs of session fidelity for the multiplicative fragment.
- A locally nameless representation of binding, eliminating the need for reasoning up to  $\alpha$ -equivalence and simplifying substitution reasoning.

*Scope of the formalization.* This thesis focuses on a mechanized reconstruction of the  $\pi$ LL system of [Montesi and Peressotti 2021], with particular emphasis on its multiplicative fragment. The

syntactic and typing components of  $\pi\text{LL}$  have been fully mechanized, including processes, session types, typing judgments, and substitution operations. The operational semantics are defined for the multiplicative fragment, with the exception of the *res*-rule.

Consequently, metatheoretic results concerning typing apply to the full system, while results involving operational semantics are restricted to the multiplicative fragment. The aim of this work is not to extend the theory, but to provide a mechanically verified account of an existing theoretical framework.

## 2 Background

This section begins by introducing the theoretical foundations underlying the formalisation following the presentation of [Montesi and Peressotti 2021]. We cover the full  $\pi\text{LL}$  calculus, including processes, types, typing, environments / hyperenvironments, and judgements, then restrict attention to the multiplicative fragment  $\pi\text{MLL}$  when discussing operational semantics.

Following [Montesi and Peressotti 2021], we use *blue* to highlight types and *red* for process terms throughout this thesis.

### 2.1 Classical Linear Logic, CLL

Linear logic [Girard 1987] is a refinement of classical and intuitionistic logic. In contrast to these systems, where formulas can be freely duplicated or discarded, linear logic restricts structural rules such as weakening and contraction, and instead treats formulas as consumable resources. This yields a resource-sensitive interpretation of logic, where each assumption must be used in a controlled manner [Di Cosmo and Miller 2023]. In this thesis, we work in the classical presentation of linear logic, where sequents are symmetric and duality replaces the distinction between assumptions and conclusions.

In particular, conjunction and disjunction split into multiplicative and additive variants. The multiplicative fragment forms the core of the system and includes the tensor connective  $A \otimes B$  and its dual  $A \wp B$ , together with the corresponding units  $1$  and  $\perp$ . These connectives capture composition and interaction of resources.

Intuitively,  $A \otimes B$  describes the joint availability of two independent resources, while  $A \wp B$  captures their dual form of interaction. Under a process interpretation,  $\otimes$  can be read as producing a value of type  $A$  and continuing as  $B$ , whereas  $A \wp B$  corresponds to receiving a value of type  $A$  and proceeding as  $B$ . The units  $1$  and  $\perp$  represent the absence of further output and input, respectively. A simplified presentation of these rules is given below:

$$\frac{}{\vdash 1} \quad \frac{\vdash \Gamma}{\vdash \Gamma, \perp} \perp \quad \frac{\vdash \Gamma, A \quad \vdash \Delta, B}{\vdash \Gamma, \Delta, A \otimes B} \otimes \quad \frac{\vdash \Gamma, A, B}{\vdash \Gamma, A \wp B} \wp$$

Here rule  $\otimes$  presents a simplified form of the tensor rule from classical linear logic. In general, the tensor connective  $A \otimes B$  combines two independent derivations. It requires a proof of  $A$  and a proof of  $B$ , each obtained from disjoint contexts, and produces a proof of the composite resource  $A \otimes B$ . Intuitively, this corresponds to combining two independent resources into a single composite resource, and can be read as producing a value of type  $A$  and then continuing as  $B$ .

The unit rule *rule 1* introduces the unit  $1$ , which represents the absence of further obligations. It can be understood as an “empty output”, corresponding to a computation that performs no further interaction, and serves as the neutral element of tensor.

Dually, rules  $\perp$  and  $\wp$  describe the behavior of the connective  $A \wp B$  and its unit  $\perp$ . The  $\wp$  rule combines resources within a single context, reflecting a form of interaction where both components

are made available to the surrounding context. From an operational perspective, this corresponds to interacting with an input of type  $A$  while continuing as  $B$ . The unit  $\perp$  represents the absence of further input, acting as an “empty input” and serving as the dual neutral element.

For example, two independent instances of the unit  $\perp$  can be introduced by the [rule 1](#) and combined via [rule  \$\otimes\$](#)  to derive  $\perp \otimes \perp$ :

$$\frac{\frac{}{\vdash \perp} \quad \frac{}{\vdash \perp}}{\vdash \perp \otimes \perp} \otimes$$

Here each leaf introduces one resource independently, and the tensor rule combines them into a single compound resource  $\perp \otimes \perp$ . After this step, the individual resources are no longer available separately, reflecting the resource-sensitive nature of the system.

This resource-oriented view of logic provides the foundation for its applications in computer science, particularly in the study of concurrency and session-typed communication. In the following sections, this perspective is used to motivate the process interpretation underlying  $\pi$ LL.

## 2.2 The $\pi$ -Calculus: Processes and Communication

The  $\pi$ -calculus [[Milner et al. 1992](#)] is a process calculus for modeling concurrent computation with dynamic communication topology, where channel names themselves can be transmitted, allowing the communication network to evolve during execution. We follow the presentation of [Montesi and Peressotti \[2021\]](#), which adopts [Wadler](#)’s convention of using square brackets ‘[-]’ for output and round parentheses ‘(-)’ for input.

Programs in  $\pi$ MLL are processes  $(P, Q, R, \dots)$ , which communicate over names  $(x, y, z, \dots)$  representing session endpoints. Sessions are binary (they involve exactly two endpoints), a discipline enforced by the typing system introduced in 2.3. We assume an infinite set of names with decidable equality. The process terms of  $\pi$ MLL are defined by the following grammar:

|            |                                                    |                                                                                                         |
|------------|----------------------------------------------------|---------------------------------------------------------------------------------------------------------|
| $P, Q ::=$ | $x[y].P$                                           | <i>output <math>y</math> on <math>x</math> and continue as <math>P</math></i>                           |
|            | $  \quad x(y).P$                                   | <i>input <math>y</math> on <math>x</math> and continue as <math>P</math></i>                            |
|            | $  \quad x[].P$                                    | <i>output (empty message) on <math>x</math> and continue as <math>P</math></i>                          |
|            | $  \quad x().P$                                    | <i>input (empty message) on <math>x</math> and continue as <math>P</math></i>                           |
|            | $  \quad \nu xy.P$                                 | <i>name restriction (cut)</i>                                                                           |
|            | $  \quad P \mid Q$                                 | <i>parallel composition</i>                                                                             |
|            | $  \quad x[L].P$                                   | <i>select left on <math>x</math> and continue as <math>P</math></i>                                     |
|            | $  \quad x[R].P$                                   | <i>select right on <math>x</math> and continue as <math>P</math></i>                                    |
|            | $  \quad x.\text{case}\{L:R, R:Q\}$                | <i>offer a binary choice between <math>P</math> or <math>Q</math> on <math>x</math></i>                 |
|            | $  \quad x[A].P$                                   | <i>output type <math>A</math> on <math>x</math> and continue as <math>P</math></i>                      |
|            | $  \quad x(X).P$                                   | <i>input a type as <math>X</math> on <math>x</math> and continue as <math>P</math></i>                  |
|            | $  \quad !x\langle z_1, \dots, z_n \rangle. \{P\}$ | <i>server (replicable process) on <math>x</math> depending on servers on <math>z_1 \dots z_n</math></i> |
|            | $  \quad x[\text{USE}].P$                          | <i>consume a server on <math>x</math> and continue as <math>P</math></i>                                |
|            | $  \quad x[\text{DUP}](x').P$                      | <i>duplicate server <math>x</math> as <math>x'</math> in <math>P</math></i>                             |
|            | $  \quad x[\text{DISP}].P$                         | <i>dispose of a server on <math>x</math></i>                                                            |
|            | $  \quad x \leftrightarrow y$                      | <i>link <math>x</math> and <math>y</math></i>                                                           |
|            | $  \quad \mathbf{0}$                               | <i>terminated process</i>                                                                               |

Term  $x[y].P$  allocates a fresh name  $y$ , outputs it over  $x$ , and proceeds as  $P$ . Dually,  $x(y).P$  inputs a name over  $x$ , binding it to  $y$  in  $P$ . Terms  $x[].P$  and  $x().P$  model empty output and input. Restriction

$\nu xyP$  connects the two endpoints  $x$  and  $y$  of  $P$  to form a session, where the order of  $x$  and  $y$  is irrelevant and  $\nu xyP = \nu yxP$ . Parallel composition  $P \mid Q$  runs  $P$  and  $Q$  concurrently, and  $0$  is the terminated process, acting as the unit of parallel composition.

*Example 2.1.* Consider the process  $P \triangleq \nu xz(x[].\mathbf{0} \mid z().y[].\mathbf{0})$ . This process creates a private session between the restricted endpoints  $x$  and  $z$ . The left process,  $x[].\mathbf{0}$ , signals termination on  $x$  before terminating, while the right process,  $z().y[].\mathbf{0}$ , waits for a termination signal on  $z$  before proceeding by sending a termination signal on  $y$  and then terminating.

This calculus provides the operational foundation for  $\pi\text{LL}$ . In the following section, we introduce the type system that governs how names are used, ensuring that each session endpoint is used exactly according to its protocol.

### 2.3 Types and Propositions-as-Sessions Correspondence

In  $\pi\text{LL}$ , session types are linear logic propositions [Caires and Pfenning 2010; Montesi and Peressotti 2021; Wadler 2014]. Each type describes the communication protocol of a channel endpoint, and duality, the key structural feature of classical linear logic, ensures that the two endpoints of every session implement complementary protocols.

The correspondence between linear logic and session types as presented in [Montesi and Peressotti 2021] following [Carbone et al. 2016; Wadler 2014] is captured directly by the type grammar. Each connective denotes a communication action, where  $A \otimes B$  describes sending a value of type  $A$  and continuing as  $B$ , while its dual  $A \wp B$  describes receiving. The units  $1$  and  $\perp$ , again, represent the absence of out and input, respectively. The additive connectives  $A \oplus B$  and  $A \& B$  capture choice in the sense of selection and offering of alternatives, while the exponentials  $!A$  and  $?A$  govern replication. The  $\forall X.A$  and  $\exists X.A$  allow polymorphic session behavior. The type grammar of  $\pi\text{LL}$  is defined as follows:

|            |               |                                  |  |               |                              |
|------------|---------------|----------------------------------|--|---------------|------------------------------|
| $A, B ::=$ | $A \otimes B$ | send $A$ , proceed as $B$        |  | $A \wp B$     | receive $A$ , proceed as $B$ |
|            | $1$           | empty output, unit for $\otimes$ |  | $\perp$       | empty input, unit for $\wp$  |
|            | $A \& B$      | offer $A$ or $B$                 |  | $A \oplus B$  | select $A$ or $B$            |
|            | $\exists X.A$ | existential, type output         |  | $\forall X.A$ | universal, type input        |
|            | $?A$          | client request                   |  | $!A$          | server accept                |
|            | $X$           | type variable                    |  | $\perp X$     | dual of type variable        |

Duality is defined structurally on session types by exchanging each connective with its dual and forms an involution, i.e.  $(A^\perp)^\perp = A$ . Consequently, the endpoints of a well-typed session are always assigned dual types, ensuring they are compatibility:

$$\begin{aligned}
 (A \otimes B)^\perp &= A^\perp \wp B^\perp & 1^\perp &= \perp & (A \oplus B)^\perp &= A^\perp \& B^\perp \\
 (!A)^\perp &= ?A^\perp & (\exists X.A)^\perp &= \forall X.A^\perp & (X)^\perp &= X^\perp
 \end{aligned}$$

### 2.4 Environments, Hyperenvironments and Judgements

Typing in  $\pi\text{LL}$  is organised into two layers. Environments track how individual names are typed, while hyperenvironments track which names are used independently in parallel.

Following the presentation in [Montesi and Peressotti 2021], *environments*  $(\Gamma, \Delta, \Xi, \dots)$  are defined as lists allowing exchange, meaning order is irrelevant. They can be written as  $\Gamma = x_1 : A, \dots, x_n : A_n$ , where each  $x_i$  is associated to its respective  $A_i$ , for  $1 \leq i \leq n$ . Environments can be merged as long as they do not share any names for example assume  $\Gamma$  and  $\Delta$  do not share any names, then  $\Gamma, \Delta$  is the

environment containing exactly the associations in  $\Gamma$  and  $\Delta$  regardless of ordering. Environments act as a partial commutative monoid, using “,” as the sum and the empty environment  $\bullet$  as unit:

$$\Gamma, \bullet = \Gamma \quad \Gamma, \Delta = \Delta, \Gamma \quad (\Gamma, \Delta), \Xi = \Gamma, (\Delta, \Xi)$$

Environments dictate how names are used, but it does not distinguish whether names are used independently or in parallel. That role is handled by *hyperenvironments*  $(\mathcal{G}, \mathcal{H}, \mathcal{I}, \dots)$ , which are collections of environments joined using the parallel operator ‘|’. They can be written as follows  $\mathcal{G} = \Gamma_1, \dots, \Gamma_n$ , where  $\Gamma_1, \dots, \Gamma_n$  is a collection of environments. Given  $\mathcal{G}$  and  $\mathcal{H}$  having no names in common then  $\mathcal{G} | \mathcal{H}$  is the hyperenvironment containing exactly the environments of  $\mathcal{G}$  and  $\mathcal{H}$ , ignoring order. Like environments, all names in a hyperenvironment are unique, they can be combined as long as they do not share any names, and they act as a partial commutative monoid using ‘|’ the sum and  $\emptyset$  as unit, where  $\emptyset$  is the hyperenvironment containing no environments:

$$\mathcal{G} | \emptyset = \mathcal{G} \quad \mathcal{G} | \mathcal{H} = \mathcal{H} | \mathcal{G} \quad (\mathcal{G} | \mathcal{H}) | \mathcal{I} = \mathcal{G} | (\mathcal{H} | \mathcal{I})$$

*Judgements*, written as  $\vdash P :: \mathcal{G}$ , states that the process  $P$  uses its free names in accordance with  $\mathcal{G}$  and can be derived using the inference rules displayed in Figure 1.

*Typing Derivations*  $(\mathcal{D}, \mathcal{E}, \mathcal{F}, \dots)$  are written as  $\vdash P :: \mathcal{G}$  and are read as the derivation  $\mathcal{D}$  having the judgement  $\vdash P :: \mathcal{G}$  as its conclusion. The meaning of this statement is that the process,  $P$ , appearing in the judgement concluding a derivation is well-typed.

---

*Multiplicative Fragment*

---

$$\begin{array}{c} \frac{\vdash P :: \mathcal{G} \mid \Gamma, x : A \mid \Delta, y : A^\perp}{\vdash \nu xy P :: \mathcal{G} \mid \Gamma, \Delta} \text{CUT} \quad \frac{\vdash P :: \mathcal{G} \quad \vdash Q :: \mathcal{H}}{\vdash P | Q :: \mathcal{G} | \mathcal{H}} \text{MIX} \quad \frac{}{\vdash 0 :: \emptyset} \text{MIX}_0 \\ \frac{\vdash P :: \Gamma, y : A \mid \Delta, x : B}{\vdash x[y].P :: \Gamma, \Delta, x : A \otimes B} \otimes \quad \frac{\vdash P :: \emptyset}{\vdash x[].P :: x : 1} 1 \quad \frac{\vdash P :: \Gamma, y : A, x : B}{\vdash x(y).P :: \Gamma, x : A \wp B} \wp \quad \frac{\vdash P :: \Gamma}{\vdash x().P :: \Gamma, x : \perp} \perp \end{array}$$

---

*Additional rules for Additives, Quantifiers and Exponentials*

---

$$\begin{array}{c} \frac{\vdash P :: \Gamma, x : A}{\vdash x[L].P :: \Gamma, x : A \oplus B} \oplus_1 \quad \frac{\vdash P :: \Gamma, x : B}{\vdash x[R].P :: \Gamma, x : A \oplus B} \oplus_2 \quad \frac{\vdash P :: \Gamma, x : A \quad \vdash Q :: \Gamma, x : B}{\vdash x.\text{case}\{L:P, R:Q\} :: \Gamma, x : A \& B} \& \\ \frac{}{\vdash x \leftrightarrow y :: x : A^\perp, y : A} \text{AX} \quad \frac{\vdash P :: \Gamma, x : A}{\vdash x[\text{USE}].P :: \Gamma, x : ?A} ? \quad \frac{\vdash P :: \Gamma}{\vdash x[\text{DISP}].P :: \Gamma, x : ?A} \text{w} \quad \frac{\vdash P :: \Gamma, x : ?A, x' : ?A}{\vdash x[\text{DUP}](x').P :: \Gamma, x : ?A} \text{c} \\ \frac{\vdash P :: z_1 : ?B_1, \dots, z_n : ?B_n, x : A}{\vdash !x(z_1, \dots, z_n). \{P\} :: z_1 : ?B_1, \dots, z_n : ?B_n, x : !A} ! \quad \frac{\vdash P :: \Gamma, x : B\{A/X\}}{\vdash x[A].P :: \Gamma, x : \exists X.B} \exists \quad \frac{\vdash P :: \Gamma, x : B \quad X \notin \text{ft}(\Gamma)}{\vdash x(X).P :: \Gamma, x : \forall X.B} \forall \end{array}$$

Fig. 1.  $\pi$ LL, typing rules.

The multiplicative rules form the core of the system. **Rule CUT** composes two processes in  $P$  sharing dual endpoints  $x : A$  and  $y : A^\perp$ , hiding them via restriction. This is the logical equivalent of communication. **Rule MIX** and **rule MIX<sub>0</sub>** compose independent processes in parallel and introduce the terminated process, respectively.

**Rule  $\otimes$**  and **rule  $\wp$**  type output and input. Note that **rule  $\otimes$**  types the sent name  $y$  in a *separate* environment from the continuation, reflecting that the two are independent resources. By contrast, **rule  $\wp$**  keeps  $y$  and the continuation  $x$  in the *same* environment, reflecting their sequential

dependency. The units **rule 1** and **rule  $\perp$**  type empty output and empty input, where **rule 1** requires the continuation to be well-typed in the empty hyperenvironment, meaning  $P$  is necessarily semantically equivalent to  $0$ .

The additive rules (**rule  $\oplus_1$** , **rule  $\oplus_2$** , and **rule  $\&$** ) type selection and offering of alternative processing paths. Note that **rule  $\&$**  requires both branches to be typed in the *same* context  $\Gamma$ , reflecting that the choice of branch is made by the other endpoint.

The exponential rules type server replication. The **rule !** requires all free names except  $x$  to be typed using  $?$ , enforcing that a server only depends on other servers. The **rule  $?$** , **rule  $w$** , and **rule  $c$**  allow a client to use, discard, or duplicate a server respectively.

The **rule  $ax$**  types a link between two dual endpoints,  $x \leftrightarrow y$ , forwarding all messages sent on  $x$  safely to  $y$  and vice versa. **Rule  $\exists$**  and **rule  $\forall$**  type polymorphic communication, where the side condition  $X \notin \text{ft}(\Gamma)$  in **rule  $\forall$** , ensures the introduced type variable,  $X$ , does not accidentally shadow one already existing in  $\Gamma$ .

*Example 2.2.* The process from [Example 2.1](#) has the typing  $\vdash \mathbf{v}zx(z().y[].0 \mid x[].0) :: y:1$ , as shown by the derivation below:

$$\begin{array}{c}
 \frac{}{\vdash 0 :: \emptyset} \text{MIX}_0 \quad \frac{}{\vdash y[].0 :: y:1} 1 \\
 \frac{}{\vdash x[].0 :: x:1} 1 \quad \frac{}{\vdash z().y[].0 :: y:1, z:\perp} \perp \\
 \frac{}{\vdash x[].0 \mid z().y[].0 :: x:1 \mid y:1, z:\perp} \text{MIX} \\
 \frac{}{\vdash \mathbf{v}zx(x[].0 \mid z().y[].0) :: y:1} \text{CUT}
 \end{array}$$


Note that, for the typing context  $x:1 \mid y:1, z:\perp$  to align with the premise of **rule CUT**, we utilise the monoidal properties of environments and hyperenvironments. In particular, using their identity laws, we instantiate the meta-variables of the rule as  $\mathcal{G} = \emptyset$  and  $\Gamma = \bullet$ . Under these instantiations, the premise of **rule CUT**:  $\vdash P :: \mathcal{G} \mid \Gamma, x:A \mid \Delta, y:A^\perp$  can be viewed as  $x:A \mid \Delta, y:A^\perp$ . This effectively instantiates the **CUT** over the dual types  $A = 1$  and  $A^\perp = \perp$  on channels  $x$  and  $y$ , with  $\Gamma = y:1$ , while ensuring the restricted endpoints implement dual session behaviours, fulfilling the requirements for a creating a valid **CUT** (session).

## 2.5 Operational Semantics


The operational semantics of  $\pi\text{LL}$  are given as a labelled transition system (LTS) defined over *typing derivations*, in which processes evolve by performing actions on their free names. The semantics follow a Structural Operational Semantics (SOS) style presentation and form the basis for the metatheoretic results of the calculus. We will be focusing on the multiplicative fragment for this part, which is given in [Figure 2](#). The full LTS is given in [Montesi and Peressotti \[2021\]](#).

*Remark 2.3.* The operational semantics rely explicitly on the structural equivalence (**rule  $=_\alpha$** ) to rename restricted variables via  $\alpha$ -equivalence prior to a transition. In paper-based presentations,  $\alpha$ -equivalence is often handled implicitly through Barendregt's Variable Convention. In a mechanised setting, however, reasoning over named variables requires freshness conditions, capture-avoiding substitution, and explicit renaming lemmas, introducing substantial proof overhead [\[Aydemir et al. 2005; Charguéraud 2012\]](#). As discussed in [Section 3](#), this motivates the use of a different binding representation for the syntax of  $\pi\text{LL}$ , one designed to reduce explicit  $\alpha$ -equivalence reasoning and simplify the treatment of bound variables.

Actions are defined in the set  $\text{Act} = \{x[], x(), x[y], x(y) \mid x, y \text{ names}\}$ , representing empty output, empty input, name output, and name input respectively. These come together with the silent label  $\tau$  for internal communication and the parallel composition of labels  $l \mid l'$  to form the

Fig. 2.  $\pi$ MLL, transition rules for derivations.





**Fig. 3.** An execution for the derivation in Example 2.2.

set of transition labels, defined as  $Lbl = \{\tau, l, (l|l') \mid l, l' \in Act, i(l) \cap i(l') = \emptyset\}$ . The restriction imposed on the parallel composition of labels exists to ensure two labels do not introduce the same name, which would break linear resource management.

*Example 2.4.* Figure 3 shows a possible execution for the derivation of the process from Example 2.1.

Each logical rule (**1**, **1 $\perp$** , **⊗**, **⊗ $\wp$** ) has a corresponding transition axiom. These follow the prefix-continuation pattern  $\pi.P \xrightarrow{\pi} P$ , where performing the action  $\pi$  leads to the remainder of the process. Rule **SYN** allows for the pairing of concurrent actions via  $l|l'$ , which is used to facilitate communication by letting two endpoints connected by a restriction perform compatible actions. When two compatible actions are performed over endpoints connected by a restriction, rules **1 $\perp$**  and **⊗ $\wp$**  simplify the term into an internal  $\tau$ -transition, modeling the synchronisation of the session.

## 2.6 Process Semantics and Erasure

While the typing derivations of  $\pi$ LL dictate the valid interactions within a session, the underlying process terms possess a standard, untyped operational behavior that closely mirrors the classic  $\pi$ -calculus. To formally relate the typed derivations to their untyped execution, we follow [Montesi and Peressotti 2021] and define an erasure function,  $proc : Der \rightarrow Proc$ , which strips away the typing information from a derivation. Recursively,  $proc$  maps the head of a typing derivation to the specific process constructor it introduces, continuing structurally onto the process's continuation. For any well-typed  $\pi$ LL term,  $proc(\mathcal{D})$  effectively extracts the pure process term from the derivation's conclusion.

*Example 2.5.* Consider the well-typed process  $P$  formed by adding a third parallel component to the derivation of Example 2.2 before the cut (A derivation of which can be seen on Figure 4):

$$\begin{array}{c}
\frac{\frac{\frac{}{\vdash \mathbf{0} :: \emptyset} \text{MIX}_0}{\vdash v[].\mathbf{0} :: v:1} \text{1} \quad \frac{\frac{\frac{}{\vdash \mathbf{0} :: \emptyset} \text{MIX}_0}{\vdash x[].\mathbf{0} :: x:1} \text{1} \quad \frac{\frac{}{\vdash \mathbf{0} :: \emptyset} \text{MIX}_0}{\vdash y[].\mathbf{0} :: y:1} \text{1}}{\vdash z().y[].\mathbf{0} :: y:1, z:\perp} \perp \\
\frac{\vdash w().v[].\mathbf{0} :: v:1, w:\perp}{\vdash (w().v[].\mathbf{0} \mid x[].\mathbf{0} \mid z().y[].\mathbf{0}) :: v:1, w:\perp \mid x:1 \mid y:1, z:\perp} \text{MIX} \\
\frac{\vdash (w().v[].\mathbf{0} \mid x[].\mathbf{0} \mid z().y[].\mathbf{0}) :: v:1, w:\perp \mid x:1 \mid y:1, z:\perp}{\vdash \nu xz (w().v[].\mathbf{0} \mid x[].\mathbf{0} \mid z().y[].\mathbf{0}) :: v:1, w:\perp \mid y:1} \text{CUT}
\end{array}$$

Fig. 4. A derivation for the process  $\nu xz w().y[].\mathbf{0} \mid x[].\mathbf{0} \mid z().y[].\mathbf{0}$

$$P = \nu xz (w().v[].\mathbf{0} \mid x[].\mathbf{0} \mid z().y[].\mathbf{0})$$

The component containing  $w$  and  $v$  is independent from the component containing  $x$ ,  $z$ , and  $y$ . This allows their respective actions to interleave freely, resulting in the following execution traces, illustrating the concurrency of the calculus:

$$\begin{array}{lll}
P \xrightarrow{\tau} w() \xrightarrow{v[]} y[] \xrightarrow{} \mathbf{0} & P \xrightarrow{\tau} w() \xrightarrow{y[]} v[] \xrightarrow{} \mathbf{0} & P \xrightarrow{\tau} y[] \xrightarrow{w()} v[] \xrightarrow{} \mathbf{0} \\
P \xrightarrow{w()} v[] \xrightarrow{\tau} y[] \xrightarrow{} \mathbf{0} & P \xrightarrow{w()} \tau \xrightarrow{v[]} y[] \xrightarrow{} \mathbf{0} & P \xrightarrow{w()} \tau \xrightarrow{y[]} v[] \xrightarrow{} \mathbf{0}
\end{array}$$

The semantics for process terms must inherently agree with the semantics of derivations for well-typed terms. As formulated in [Montesi and Peressotti 2021], this relationship is established functorially. Let  $\mathcal{P}(X) = \{X' \mid X' \subseteq X\}$  and  $\mathcal{L}(X) = X \times \text{Lbl}$  be endofunctors representing the states and labeled transitions, respectively. The erasure function *proc* acts as a homomorphism from LTS of derivations ( $\text{LTS}_d$ ) to the LTS of processes ( $\text{LTS}_p$ ), ensuring that the square of these LTS mappings commutes.

In the context of operational semantics, this functorial homomorphism yields an Erasure theorem. It guarantees that the Labeled Transition System for derivations and the SOS for untyped processes simulate each other exactly:

**THEOREM 2.6 (ERASURE).** *The process LTS ( $\text{LTS}_p$ ) enjoys erasure with respect to the derivation LTS ( $\text{LTS}_d$ ). Concretely, for any valid typing derivation  $\mathcal{D}$  and label  $l \in \text{Lbl}$ :*

- (1) **Soundness of Erasure:** *If  $\mathcal{D} \xrightarrow{l} \mathcal{D}'$ , then  $\text{proc}(\mathcal{D}) \xrightarrow{l} \text{proc}(\mathcal{D}')$ .*
- (2) **Completeness of Erasure:** *If  $\text{proc}(\mathcal{D}) \xrightarrow{l} P'$ , then there exists a derivation  $\mathcal{D}'$  such that  $\mathcal{D} \xrightarrow{l} \mathcal{D}'$  and  $\text{proc}(\mathcal{D}') = P'$ .*

**COROLLARY 2.7 (TYPABILITY PRESERVATION).** *If  $P$  is well-typed and  $P \xrightarrow{l} P'$ , then  $P'$  is well-typed.*

This theorem allows us to define the operational target for  $\pi$ LL processes without the syntactic overhead of typing environments. Using the erasure formulation, we can transform the operational semantics for derivations directly into an SOS specification for processes.

**Remark 2.8.** The untyped LTS for processes relies heavily on explicit side conditions regarding free and introduced names to prevent the collision of bound variables (e.g., ensuring  $i(l) \cap i(l') = \emptyset$ ). However, when operating at the level of typing derivations, these side conditions become inherently redundant. Because  $\pi$ LL enforces linear resource management, the structural rules of the typed LTS dictate that hyperenvironments can only be merged when their domains are strictly disjoint. Consequently, any well-typed derivation inherently satisfies the name-distinctness and freshness

$$\begin{array}{c}
\begin{array}{c}
x[x'].P \xrightarrow{x[x']} P \quad x(x').P \xrightarrow{x(x')} P \quad x[].P \xrightarrow{x[]} P \quad x().P \xrightarrow{x()} P \\
\frac{P \xrightarrow{l} P' \quad i(l) \cap f(Q) = \emptyset}{P \mid Q \xrightarrow{l} P' \mid Q} \text{PAR}_1 \quad \frac{Q \xrightarrow{l} Q' \quad i(l) \cap f(P) = \emptyset}{P \mid Q \xrightarrow{l} P \mid Q'} \text{PAR}_2 \\
\frac{P \xrightarrow{l} P' \quad Q \xrightarrow{l'} Q' \quad i(l \parallel l') \cap f(P \mid Q) = \emptyset}{P \mid Q \xrightarrow{l \parallel l'} P' \mid Q'} \text{SYN} \quad \frac{P \xrightarrow{x[x'] \parallel y(y')} P'}{vxy P \xrightarrow{\tau} vxy vx'y' P'} \otimes \otimes \\
\frac{P \xrightarrow{x[] \parallel y()} P'}{vxy P \xrightarrow{\tau} P'} 1\bot \quad \frac{P \xrightarrow{l} P' \quad x, y \notin f(l) \cup i(l)}{vxy P \xrightarrow{l} vxy P'} \text{RES} \quad \frac{P =_{\alpha} Q \quad Q \xrightarrow{l} Q'}{P \xrightarrow{l} Q'} =_{\alpha}
\end{array}
\end{array}$$

$$\begin{array}{c}
i(l \parallel l') = i(l) \cup i(l') \\
f(l \parallel l') = f(l) \cup f(l') \\
i(\tau) = \emptyset \\
f(\tau) = \emptyset \\
i(x[]) = i(x()) = \emptyset \\
f(x[]) = f(x()) = \{x\} \\
i(x[x']) = i(x(x')) = \{x'\} \\
f(x[x']) = f(x(x')) = \{x\}
\end{array}$$

Fig. 5.  $\pi$ MLL, transition rules for processes.

requirements of the underlying untyped calculus, shifting the burden of safety from dynamic transition checks to static structural typing.

**PROPOSITION 2.9.** *The operational semantics of  $\pi$ LL processes strictly respects the introduction of names via transition labels. Formally, if  $P \xrightarrow{l} P'$ , then  $i(l) \cap f(P) = \emptyset$ , and If  $P$  is well-typed, then  $i(l) = f(P') \setminus f(P)$  (TODO: Evaluate relevance - delete depending on if this property is mentioend in later proofs).*

**Remark 2.10.** **Proposition 2.9** is vital for mechanising the metatheory of  $\pi$ LL. Because proof assistants tools require explicit tracking of fresh and bound variable to prevent accidental capture, this proposition acts as the formal bridge between the dynamic semantics and the static typing. It guarantees that a well-typed process evolves predictably, ensuring that it does not spontaneously generate resources during a transition. (TODO: same as **Proposition 2.9**)

## 2.7 Semantics of Typing Environments

Just as processes evolve through computation, the typing environments that govern them must also evolve to accurately track the state of a session. In  $\pi$ LL, types capture communication protocols. Consequently, the transition rules for typing environments specify exactly which actions are permitted and how the available resources are consumed by those actions.

Similar to how we defined the process erasure function, we again follow **Montesi and Peressotti [2021]** and define a function  $env : Der \rightarrow Env$  that maps a valid typing derivation directly to the hyperenvironment in its conclusion (i.e.,  $env(\vdash P :: \mathcal{G}) = \mathcal{G}$ ). Using this extraction, **Montesi and Peressotti [2021]** derive a LTS specifically for typing environments ( $lts_e$ ), shown in **Figure 6**.

In this LTS, observable actions actively consume resources from the environment. For example, if a hyperenvironment contains a waiting channel  $x : \perp$ , performing an input action  $x()$  will consume this type, removing it from the resulting environment. Conversely, internal communications do not alter the available external resources. This is formalised by the  $\tau$ -reflexivity axiom,  $\mathcal{G} \mid \Gamma \xrightarrow{\tau} \mathcal{G} \mid \Gamma$ , which states that any valid hyperenvironment transitions to itself under a silent  $\tau$ -action.

The requirement that a process can only perform actions permitted by its typing environment is known as **Session Fidelity**. Just as the *proc* function established operational correspondence for processes, the *env* function establishes a correspondence for environments.

While the rules in **Figure 6** define how environments can transition in isolation, the execution of a well-typed program requires the process and its environment to step in unison. This co-dependent relationship has processes only perform actions permitted by its typing environment, while the

$$\begin{array}{c}
x:1 \xrightarrow{x[]} \emptyset \quad \Gamma, x:\perp \xrightarrow{x()} \Gamma \quad \Gamma, \Delta, x:A \otimes B \xrightarrow{x[x']} \Gamma, x':A \mid \Delta, x:B \quad \Gamma, x:A \wp B \xrightarrow{x(x')} \Gamma, x':A, x:B \\
\frac{\mathcal{G} \xrightarrow{l} \mathcal{G}'}{\mathcal{G} \mid \mathcal{H} \xrightarrow{l} \mathcal{G}' \mid \mathcal{H}} \text{PAR}_1 \quad \frac{\mathcal{H} \xrightarrow{l} \mathcal{H}'}{\mathcal{G} \mid \mathcal{H} \xrightarrow{l} \mathcal{G} \mid \mathcal{H}'} \text{PAR}_2 \quad \frac{\mathcal{G} \xrightarrow{l} \mathcal{G}' \quad \mathcal{H} \xrightarrow{l'} \mathcal{H}'}{\mathcal{G} \mid \mathcal{H} \xrightarrow{ll'} \mathcal{G}' \mid \mathcal{H}'} \text{SYN} \\
\frac{\mathcal{G} \mid \Gamma, x:A^\perp \mid \Delta, y:A \xrightarrow{l} \mathcal{G}' \mid \Gamma', x:A^\perp \mid \Delta', y:A}{\mathcal{G} \mid \Gamma, \Delta \xrightarrow{l} \mathcal{G}' \mid \Gamma', \Delta'} \text{RES} \quad \frac{\mathcal{G} \mid x:1 \mid \Gamma, y:\perp \xrightarrow{x[]|y()} \mathcal{G} \mid \Gamma}{\mathcal{G} \mid \Gamma \xrightarrow{\tau} \mathcal{G} \mid \Gamma} 1_\perp \\
\frac{\mathcal{G} \mid \Gamma, \Delta, x:A \otimes B \mid \Xi, y:A^\perp \wp B^\perp \xrightarrow{x[x']|y(y')} \mathcal{G} \mid \Gamma, x:B \mid \Delta, x':A \mid \Xi, y:B^\perp, y':A^\perp}{\mathcal{G} \mid \Gamma, \Delta, \Xi \xrightarrow{\tau} \mathcal{G} \mid \Gamma, \Delta, \Xi} \otimes \wp
\end{array}$$

Fig. 6.  $\pi$ MLL, transition rules for typing environments.

environment dynamically mirrors the physical control flow of the process, and is formally captured by the property of **Session Fidelity**. Just as the *proc* function established operational correspondence for the untyped calculus, the *env* function establishes this exact behavioral correspondence for the typing contexts.

**THEOREM 2.11 (SESSION FIDELITY).** *Let  $\mathcal{D}$  be a valid typing derivation such that  $\text{env}(\mathcal{D}) = \mathcal{G}$  and  $\text{proc}(\mathcal{D}) = P$ . If the process transitions such that  $P \xrightarrow{l} P'$ , then the environment can mimic this transition*

$$\mathcal{G} \xrightarrow{l} \mathcal{G}',$$

*and there exists a new valid derivation  $\mathcal{D}'$  for  $P'$  under  $\mathcal{G}'$ .*

To illustrate Session Fidelity in action, we can observe how the environment is constrained to follow the exact execution traces of the process it types.

**Example 2.12 (Session Fidelity in Execution).** Recall the process derivation from [Example 2.5](#), typed by the environment  $\mathcal{G} = v:1, w:\perp \mid y:1$ . The process begins by performing an observable input action on  $w$ . By Session Fidelity, the environment must faithfully mimic this action, consuming the corresponding type:

$$v:1, w:\perp \mid y:1 \xrightarrow{w()} v:1 \mid y:1$$

The resulting derivative of the process is now typed by this new environment. At this intermediate stage, if we were to view the environment LTS in isolation, the pure semantics would theoretically allow an output on  $y$  to fire immediately, as  $y:1$  is an available, unblocked resource.

However, Session Fidelity dictates a strict co-dependence. The environment merely mirrors the process, and the underlying process term strictly guards the  $y[]$  action behind a prefix, requiring a prior  $\tau$ -synchronisation. Therefore, the environment must wait. When the process finally resolves its internal communication, the environment mimics it via the  $\tau$ -reflexivity axiom, remaining structurally unchanged:

$$v:1 \mid y:1 \xrightarrow{\tau} v:1 \mid y:1$$

Only after this synchronisation unblocks the continuation can the process perform the  $y[]$  action, which the environment once again mirrors by consuming the final resource:

$$v:1 \mid y:1 \xrightarrow{y[]} v:1$$

**Remark 2.13 (The Mechanisation Gap).** Session Fidelity is the ultimate touchstone for validating the safety of a session-typed language. However, verifying this property in a proof assistant exposes

a stark gap between paper mathematics and mechanised logic. On paper, the invariant that an environment *dynamically mirrors* a process is highly intuitive, and transitions are easily massaged using structural equivalences and implicit variable conventions. In a machine, however, formally proving that types are consumed correctly requires a rigorous, algorithmic representation of bound variables and explicit freshness. The theoretical elegance of the environment LTS cannot be mechanised without first establishing a concrete architectural foundation for these variables, which dictates the design choices discussed in the following chapters.

## 2.8 Managing Bound Variables

Having established the operational semantics of  $\pi$ LL and the ultimate goal of verifying Session Fidelity, we must address the treatment of bound variables and freshness, a central technical challenge that underpins the entire formalisation.

As we have seen, many constructs in  $\pi$ LL introduce scoped identifiers. Restriction binds pairs of session endpoints, input constructs bind received channel names, polymorphic communication binds type variables, and duplication constructs introduce fresh channel endpoints. These bindings appear throughout both the process syntax and the labelled transition system.

The concrete choice of bound and introduced names is semantically irrelevant provided the binding structure is preserved. For example, renaming a restricted channel does not change the behaviour of a process. However, reasoning formally about substitutions, transitions, and typing derivations requires these names, scopes, and freshness conditions to be tracked explicitly.

This principle is formalised as  $\alpha$ -equivalence, which identifies terms that differ only in the names of their bound variables. For instance, the following two processes are  $\alpha$ -equivalent:

$$vab(a[].\mathbf{0} \mid b().x[].\mathbf{0}) =_{\alpha} vzw(z[].\mathbf{0} \mid w().x[].\mathbf{0})$$

The operational semantics of  $\pi$ LL relies on  $\alpha$ -equivalence using  $\text{rule } =_{\alpha}$  to rename bound variables on the fly, ensuring that restricted scopes and input prefixes do not accidentally capture free names as processes evolve.

Substitution raises this exact issue. Consider the substitution where the free name  $x$  is replaced by  $z$  in the right-hand process from the previous example:

$$vzw(z[].\mathbf{0} \mid w().x[].\mathbf{0})$$

Here  $z$  is already bound by the restriction, so naively replacing  $x$  with  $z$  would cause the restriction to capture the newly inserted  $z$ . Before the substitution can proceed, and because restriction binds session endpoints as pairs, both the bound  $z$  and  $w$  must be  $\alpha$ -renamed to fresh names, such as to  $a$  and  $b$ . Similarly, any binder within the process that would capture the substituted name must also be renamed before substitution continues.

In paper-based presentations this is commonly handled by *Barendregt's variable convention*, which assumes bound variables are always chosen distinct from free variables and one another whenever convenient. This makes proofs readable but remains an informal meta-level assumption.

A proof assistant requires every freshness condition and renaming step to be represented and verified explicitly [Aydemir et al. 2005]. A direct mechanisation using named variables therefore introduces substantial proof overhead for capture avoidance and  $\alpha$ -equivalence reasoning. To avoid this, the formalisation adopts a *locally nameless* representation of binding, which eliminates  $\alpha$ -equivalence as a proof obligation by construction. This is the subject of the following chapter.

## 3 Binding Architecture

This chapter develops the locally nameless binding architecture used throughout the formalisation. We begin by stating the design goals that motivate the approach (Section 3.1), then introduce the

locally nameless representation and its core operations (Sections 3.2 and 3.3), and conclude with co-finite quantification (Section 3.4), which governs how inductive definitions interact with binders.

### 3.1 Design Goals and the Approach

This chapter introduces the binding representation used throughout the formalisation. The representation is motivated by three primary design goals:

- (1) **Eliminate  $\alpha$ -equivalence overhead.**  $\alpha$ -equivalent terms should be represented identically, reducing  $\alpha$ -equivalence reasoning to definitional equality within the proof assistant.
- (2) **Preserve human readability.** Terms should remain explicit and concrete, keeping substitution and freshness reasoning close to their paper-based presentation.
- (3) **Avoid informal meta-level conventions.** The representation should not rely on Barendregt’s convention or unverified freshness assumptions; capture avoidance and freshness conditions should instead be enforced by construction.

Although both channels and type variables involve binding, they exhibit fundamentally different operational behaviour. Session channels participate in dynamic communication and substitution across concurrent processes, while polymorphic type variables remain statically scoped within typing derivations. The formalisation therefore adopts different binding strategies for these two classes of identifiers while maintaining a common underlying binding architecture.

To achieve these goals, we adopt a locally nameless binding architecture for the calculus [Charguéraud 2012], separating free identifiers from bound occurrences. Bound variables are represented using indices, causing  $\alpha$ -equivalent terms to coincide structurally, while free variables remain explicit and readable.

This separation simplifies substitution by restricting it to free variables while leaving bound indices untouched, thereby avoiding accidental capture by construction. As a result, many proofs that traditionally require substantial renaming infrastructure reduce instead to structural recursion and index manipulations.

The remainder of this chapter develops the theoretical machinery underlying this representation.

### 3.2 The Locally Nameless Representation

The locally nameless representation, as presented by Charguéraud [2012], separates variables into two disjoint syntactic categories: free names and bound indices. This combines the readability of named syntax for free variables with a structural representation of binding.

- (1) **Bound variables** are represented using *De Bruijn indices* [de Bruijn 1972], where each index denotes the distance to its corresponding binder. Since the representation is purely positional,  $\alpha$ -equivalent terms become structurally identical.
- (2) **Free variables** retain explicit names, preserving the readability of open terms and allowing standard name-based reasoning about environments and transition labels.

By structurally separating these two concepts, the locally nameless approach simplifies standard operations such as capture-avoiding substitution. Since free variables and bound variables inhabit distinct syntactic categories, substitution on free variables is defined to not interact with bound indices. We apply this theory to both the structural and operational layers of  $\pi$ LL.

*LN Syntax for Types.* In the original presentation of  $\pi$ LL, polymorphic types rely on standard named variables. Under the locally nameless architecture, we expand the syntax of type variables to explicitly distinguish between free and bound occurrences. The grammar for types is updated as

follows:

$$\begin{aligned} T &::= \text{free}(x) \mid \text{bound}(k) & (\text{where } x, k \in \mathbb{N}) \\ A, B &::= \dots \mid \exists A \mid \forall A \mid X \mid X^\perp & (\text{where } X \in T) \end{aligned}$$

Here, a free type variable carries a unique identifier  $x$ , while a bound type variable carries a De Bruijn index  $k$ . For example, the polymorphic type  $\forall X. \exists Y. (X \wp Y)$  is structurally represented without names using the new syntax as  $\forall. \exists. (\text{bound}(1) \wp \text{bound}(0))$ , where 0 points to the inner existential binder and 1 points to the outer universal binder.

*LN Syntax for Processes.* Similarly, the standard process syntax displayed in [Section 2.2](#) is updated to reflect the locally nameless separation of channel endpoints. A channel is split into its bound and free counterparts, where bound channel names are represented by indices, while free channels remain explicit identifiers.

Naively removing the bound names from the process syntax would turn binding operations such as  $x[y].P$  and  $x(y).P$  into  $x[\cdot].P$  and  $x(\cdot).P$  making them syntactically indistinguishable from their non-binding unit counterparts. To combat this we annotate anonymous binders using placeholder symbols. We use  $\bullet$  for a bound process name, and  $\diamond$  for a bound type variable.

$$\begin{aligned} \text{Channel} &::= \text{free}(x) \mid \text{bound}(k) & (\text{where } x, k \in \mathbb{N}) \\ P, Q &::= \dots \mid x[\bullet].P \mid x(\bullet).P \mid \nu(\bullet, \bullet).P \mid x(\diamond).P & (\text{where } x \in \text{Channel}) \end{aligned}$$

While the visual anonymity of these constructors is evident from the grammar, their mechanical implementation relies on precise index management. For a multi-name binder like restriction,  $\nu(\bullet, \bullet).P$ , the two dual endpoints are simultaneously mapped to distinct indices within the continuation of  $P$ , namely 0 and 1. For communication prefixes like  $x(\bullet).P$ , the payload is implicitly bound as index 0 within  $P$ , while the explicit free channel  $x$  is deliberately preserved as the subject. By maintaining explicit free names strictly for these outward-facing communications, the transition labels and operational semantics of the calculus remain completely unaffected by the architectural shift.

While the locally nameless architecture elegantly resolves  $\alpha$ -equivalence, moving between abstracted and instantiated forms requires specialised operations. To execute a substitution or evaluate a process across a binder, we must formally define how to open and close these terms.

### 3.3 Opening and Closing Binders

Because the locally nameless representation separates free and bound variables, standard substitution cannot be applied uniformly under binders. Instead, moving between bound and free representations is handled by two dual operations, namely *opening* and *closing* [[Charguéraud 2012](#)].

*Opening (Instantiation).* Opening is the operation of traversing into a binding construct and replacing the exposed bound index with a concrete free name. Formally, opening a process  $P$  at depth  $k$  with a free name  $x$  is denoted  $\{k \mapsto x\}P$ .

The operation proceeds via structural recursion. A depth counter, initially set to 0, tracks the number of binders traversed. As the recursion descends through the syntax tree, the counter increments when passing a new binder. Consider being at depth  $k$ , then traversing a single-name binder like  $x(\bullet).P$  increments the depth to  $k+1$ , while traversing a double-name binder like  $\nu(\bullet, \bullet).P$  would increment it to  $k+2$ .

When the recursion reaches a bound variable  $\text{bound}(n)$ , the index  $n$  is compared against the current depth  $k$ . If  $n = k$ , the variable refers to the binder currently being opened and is replaced by the free name  $\text{free}(x)$ . Otherwise, the index refers to a different binder and is left unchanged.



Crucially, when the opening operation encounters an existing free variable  $\text{free}(y)$ , it simply ignores it. This guarantees that opening is entirely capture-avoiding by construction. In the operational semantics, when a single binder like an input prefix  $z(\bullet).P$  receives a channel  $x$ , the continuation  $P$  is opened as  $\{0 \mapsto x\}P$ . For a restriction  $\nu(\bullet, \bullet)P$ , the dual endpoints  $x$  and  $y$  are opened sequentially as  $\{0 \mapsto x\}(\{1 \mapsto y\}P)$ .

*Closing (Abstraction).* Closing is the exact inverse of opening. It is the operation of taking an open term and abstracting a specific free name  $x$  into a bound index at depth  $k$ , denoted  $\{k \leftarrow x\}P$ . While recursing, bound indices are ignored as they cannot contain free names. When the recursion encounters a free term like  $\text{free}(y)$ , it checked if  $y = x$ . If they match, the free name is replaced by  $\text{bound}(k)$ .

This operation is primarily used when constructing new binders, such as moving a free channel into a restricted scope. To close multiple free names simultaneously, such as abstracting  $x$  and  $y$  into a restriction, the operations are applied sequentially, by first closing  $y$  at depth  $k + 1$ , followed by closing  $x$  at depth  $k$ .

*Locally Closed Terms.* The separation of names and indices introduces a new syntactic anomaly, namely that a term can contain a *dangling* index that does not point to any enclosing binder. For example, the process term  $\text{bound}(0)[].\mathbf{0}$  is syntactically valid, matching the syntax for an empty send, but mathematically it is meaningless because the index 0 has no corresponding binder.

To filter out these meaningless terms, the locally nameless architecture introduces the predicate of being *locally closed*, denoted  $lc(P)$ . A term is locally closed if and only if all of its De Bruijn indices are properly bound by an enclosing constructor.

The locally closed predicate excludes terms containing dangling indices and therefore characterises the well-formed syntax of the calculus. Establishing this predicate is not only a static validation step, but rather it forms a key invariant throughout the mechanisation, which guarantees the soundness of the calculus. Generally, any subsequent definition, such as variable substitution, structural congruence, or operational semantics, must be mathematically proven to *preserve* the locally closed property. In practice, this is achieved through the *regularity of typing*, which requires proving that all valid typing derivations inherently produce locally closed terms. Consequently, operations proven to preserve typing will naturally carry this invariant as a structural guarantee.

### 3.4 Co-finite Quantification

With the syntax partitioned into names and indices, and the mechanics of opening and closing established, the final theoretical hurdle is defining how inductive relations, such as typing judgements and labelled transitions, operate across binders.

In paper-based proofs, when a rule requires opening a binder, Barendregt’s variable convention allows the author to simply state: *pick a fresh name*  $x$ . In a proof assistant, this informal generation of a fresh name must be encoded into the constructors of the inductive relations. Two naive approaches to bind this fresh name, both of which introduce severe mechanical flaws [Aydemir et al. 2008; Charguéraud 2012]:

- (1) **Existential Quantification** ( $\exists x \notin f(P)$ ): The rule requires the property to hold for at least one fresh name. While this makes proof construction straightforward, it yields weak induction hypotheses, since induction only provides the property for one particular fresh name chosen during the original derivation.
- (2) **Universal Quantification** ( $\forall x$ ): The rule requires the property to hold for every name. This provides a strong induction hypothesis, but is too restrictive for proof construction,



since the property cannot generally be shown for names that already occur free in the surrounding context.

To solve this dilemma, the locally nameless architecture employs **co-finite quantification** [Aydemir et al. 2008; Charguéraud 2012]. Rather than quantifying over one name or all names, the inductive rules quantifies universally over all names *except* those not in an arbitrary, finite set  $L$ . Formally, a rule requiring a fresh name uses the premise:  $\forall x \notin L$ .

By deferring the exact choice of the forbidden set  $L$  to the user at the time of proof, co-finite quantification achieves the best of both worlds. When constructing a proof tree, the user can instantiate  $L$  as the set of all free variables currently in the environment, process, and transition labels ( $f(P) \cup f(\Gamma) \cup \dots$ ). This ensures the premise is trivially satisfiable. Conversely, when performing structural induction, the co-finite universal quantifier guarantees that the induction hypothesis applies to *any* fresh name the user could possibly need, provided they pick one outside the finite set  $L$ .

*Application to  $\pi$ LL Types.* Although type variables are represented using the locally nameless separation of free and bound constructors, their metatheory is handled using depth-indexed reasoning rather than cofinite quantification. Because type variables are static and do not dynamically communicate across parallel scopes like channels do, managing freshness sets  $L$  for them introduces unnecessary proof overhead.

Instead, the system evaluates polymorphic typing rules using a purely depth-indexed approach, powered by a standard De Bruijn shifting operator. We define the shifting of a type  $A$  by a distance  $d$  at a cutoff depth  $c$ , denoted mathematically as  $\uparrow_d^c A$ . Using this infrastructure, consider the typing rule for the universal type input ( $\forall.B$ ):

$$\frac{n+1 \vdash P :: \Gamma^{\uparrow^t}, x:B}{n \vdash x(\diamond).P :: \uparrow_0^1 \Gamma, x:\forall.B} \forall_{DB}$$

Rather than opening the binder with a fresh explicit name, the typing judgment simply increments a depth counter  $n+1$ , legally bringing index 0 into scope. To maintain correct index distances, the typing environment  $\Gamma$  is explicitly shifted by a distance of 1 at cutoff 0 ( $\uparrow_0^1 \Gamma$ ), incrementing all free type indices within the context to prevent accidental capture.

*Application to  $\pi$ LL Channels.* This co-finite quantification strategy fundamentally reshapes the inductive definitions for process channels. For example, consider the typing rule that introduces the prefix  $x(\bullet).P$ . Under a paper-based presentation, the premise implicitly requires a fresh session name  $y$ . Under our mechanised architecture, the rule uses co-finite quantification and the opening operation:

$$\frac{\forall y \notin L, \quad n \vdash \{0 \mapsto y\}P :: \Gamma, x:B, y:A}{n \vdash x(\bullet).P :: \Gamma, x:A \wp B} \wp_{LN}$$

Here, the rule states that the process is well-typed if, for any fresh channel variable  $y$  outside some finite set of forbidden variables  $L$ , the opened process  $P$  is well-typed under the environment extended with the opened payload. By utilizing this mixed evaluation strategy, co-finite quantification for dynamic channels, and depth-indexed evaluation for static types, the architecture minimizes bureaucratic renaming overhead while maintaining a unified underlying syntax.

With this mathematical foundation established, the conceptual theory of  $\pi$ LL is fully prepared for implementation. The following chapter details how this locally nameless architecture is concretely mechanised in the Lean 4 proof assistant.

## 4 Formalisation

This section describes the mechanisation of  $\pi$ LL in Lean 4, bridging the theoretical foundations developed in [Sections 2 and 3](#) and their concrete implementation. The formalisation is organised into three top-level namespaces: `PiLL.Model`, which defines the static structure of the calculus, `PiLL.Semantics`, which defines its operational behaviour, and `PiLL.SessionFidelity`, which establishes the metatheoretic results connecting the two. Each component is presented in dependency order, highlighting design decisions and divergences from [Montesi and Peressotti \[2021\]](#).

### 4.1 Model

The `Model` namespace covers the full  $\pi$ LL calculus and is structured around the layers introduced in [Section 2](#). `Base.lean` provides shared syntactic infrastructure used across the model. Session types are defined in `STypes/`, processes in `Processes/`, and typing contexts in `Environments/` and `HyperEnvironment/`. Typing judgements and their properties are collected in `Judgement/`, which includes inversion lemmas, linearity properties, substitution, and the co-finite freshness infrastructure described in [Section 3.4](#).

**4.1.1 Session Types.** Session types are defined in `STypes/Basic.lean` as the inductive datatype `Types`. Following the locally nameless syntax established in [Section 3.2](#), type variables are isolated into a separate inductive type, `TVar`, which explicitly distinguishes between free and bound occurrences:

```
inductive TVar : Type
  | free (x : FVar)
  | bound (x : BVar)
```

Here, `FVar` represents a free type variable identified by a natural number, while `BVar` represents a bound type variable identified by a De Bruijn index (Which is also just a natural number). Because bound variables are represented purely by positional indices,  $\alpha$ -equivalent types become definitionally equal in Lean, eliminating the need for quotient or renaming infrastructure.

The remaining connectives directly mirror the theoretical grammar of the calculus. [Table 3](#) illustrates the translation of some connectives from their paper presentation (introduced in [Section 2.3](#)) into the constructors of the inductive `Types` declaration.

**Table 3.** Comparison of standard, locally nameless, and Lean representations for  $\pi$ LL types.

| Standard      | Locally Nameless      | Lean Constructor                  | Binds  |
|---------------|-----------------------|-----------------------------------|--------|
| $A \otimes B$ | $A \otimes B$         | <code>tensor (A B : Types)</code> | None   |
| $1$           | $1$                   | <code>one</code>                  | None   |
| $\forall X.A$ | $\forall \diamond .A$ | <code>forall_ (A : Types)</code>  | 1 type |
| $\exists X.A$ | $\exists \diamond .A$ | <code>exists_ (A : Types)</code>  | 1 type |
| $X$           | $X$                   | <code>var (v : TVar)</code>       | None   |
| $X^\perp$     | $X^\perp$             | <code>varDual (v : TVar)</code>   | None   |

Notably, the polymorphic binders  $\forall.A$  and  $\exists.A$  reflect the depth-indexed approach detailed in [Section 3.4](#). The constructors keep the bound De Bruijn index implicit, taking only the continuation body  $A$  as an argument (an underscore “`_`” has been appended to the constructor names to prevent namespace collisions with Lean’s built-in `forall` and `exists` keywords).

**Duality** is implemented as a structurally recursive function `Types.dual`, mapping each constructor to its dual. For example:

```
def Types.dual : Types -> Types
| .tensor A B => .parr (dual A) (dual B)
| .forall_ A  => .exists_ (dual A)
| .one       => .bot
| ...
```

The fundamental property that duality is an involution,  $(A^\perp)^\perp = A$ , is proved in `STypes/Lemmas.lean` by structural induction. Furthermore, because duality interacts deeply with the depth-indexed binding architecture, the formalisation establishes a suite of essential commutation lemmas. Crucially, we prove that duality preserves local closure ( $A.\text{lc} \iff A^\perp.\text{lc}$ ), and that it commutes with both De Bruijn shifting ( $\uparrow_d^c(A^\perp) = (\uparrow_d^c A)^\perp$ ) and substitution ( $B^\perp\{A/k\} = (B\{A/k\})^\perp$ ). It has also been established that duality is injective ( $A^\perp = B^\perp \iff A = B$ ).

**Local closure** for types is defined in `STypes/Properties.lean` as a depth-indexed predicate `Types.lc (k : Nat) : Types → Prop`, where the depth parameter  $k$  tracks how many binders the recursion has traversed.

At the leaves of the syntax tree, local closure is evaluated by the helper predicate `TVar.lc`. Because free variables are explicit they do not contain any indices, as such they are inherently locally closed at any depth. Conversely, a bound variable is only locally closed if its De Bruijn index  $i$  is strictly less than the current depth  $k$ , ensuring it points to an enclosing scope:

```
def TVar.lc : Nat -> TVar -> Prop
| _, .free _   => True
| k, .bound i  => i < k
```

The primary `Types.lc` predicate is then defined by structural recursion. For non-binding constructors, the depth parameter distributes across the sub-terms. When a variable is encountered, the predicate delegates the check to `TVar.lc`. When descending into the body of a polymorphic binder, the depth counter increments by one, bringing the newly introduced index safely into scope:

```
def Types.lc : Nat -> Types -> Prop
| _, .one       => True
| k, .var v      => TVar.lc k v
| k, .tensor A B => A.lc k ∧ B.lc k
| k, .forall_ A  => A.lc (k + 1)
| k, .exists_ A  => A.lc (k + 1)
| ...
```

Finally, a type is defined as *closed* if it contains no dangling indices at the top level, meaning it is locally closed at depth 0. This is registered as the abbreviation `Types.lc_0`.

**Shifting** is defined in `STypes/Shift.lean` following the standard generic De Bruijn shifting operation ( $\uparrow_c^d$ ) established in [Section 3.4](#). The implementation is split into a structural traversal over types and a base-case evaluation for variables.

The operation is mapped across the `Types` level via structural recursion, utilizing a cutoff depth  $d$  and a shift amount  $c$ . For standard connectives, the shift simply distributes across sub-terms. Crucially, because the operation traverses the syntax tree top-down, whenever the recursion descends into a polymorphic binder like `forall_`, the cutoff depth  $d$  increments by one ( $d + 1$ ). This increment ensures the algorithm constantly tracks exactly how many local binders it has entered:

```

883 def Types.shift (d c : Nat) : Types -> Types
884 | .var v          => .var (v.shift d c)
885 | .tensor A B     => .tensor (A.shift d c) (B.shift d c)
886 | .forall_ A      => .forall_ (A.shift (d + 1) c)
887 | t               => t

```

When the recursion reaches the variable level, explicit free FTVar names remain unaffected. For bound BTVar indices, the function compares the De Bruijn index  $i$  against the accumulated cutoff depth  $d$ . At this leaf node,  $d$  represents the exact number of local binders traversed during the current shift. Therefore, if  $i < d$ , the index points to one of those recently traversed binders; the variable is safely bound *internally* within the term and remains unchanged. Conversely, if  $i \geq d$ , the index points to a binder that was *never* traversed, meaning it is bound *externally* in the surrounding environment. To prevent these external variables from being accidentally captured by newly inserted binders, their indices are incremented by the shift amount  $c$  to safely jump over the new additions:

```

897 def TVar.shift (d c : Nat) : TVar -> TVar
898 | .bound i => if i < d then .bound i else .bound (i + c)
899 | .free n  => .free n

```

**Substitution** is defined in STypes/Subst.lean as the eager substitution of a payload type  $A$  for a target De Bruijn index  $k$  within a host type  $B$ , denoted as  $B\{A//k\}$ .

At the variable level, the function evaluates the bound De Bruijn index  $i$  against the target depth  $k$ . If  $i = k$ , the target has been reached, and the variable is replaced by the instantiated type  $A$ . If  $i < k$ , the variable is bound by an inner scope and is unaffected. Finally, if  $i > k$ , the variable points outward past the current term. Then since instantiating the index  $k$  effectively removes that index and associated binder leaving a whole, therefore the indices pointing outward must be decremented by one ( $i - 1$ ) to close the gap to maintain correct distanceing to the outer environment:

```

909 | .var (.bound i) =>
910   if i == k then A
911   else if i > k then .var (.bound (i - 1))
912   else .var (.bound i)

```

This operation is lifted to the Types level via structural recursion. For standard connectives, the substitution simply distributes across the sub-terms. However, when descending into a polymorphic binder such as forall\_, two critical protective adjustments must occur. First, the target depth  $k$  increments by one ( $k + 1$ ) to account for crossing the new binder. Second, the payload type  $A$  is explicitly shifted (shift 0 1  $A$ ) to ensure that any free variables residing inside  $A$  are not accidentally captured by this newly entered scope:

```

920 | .forall_ B => .forall_ (B.subst (shift 0 1 A) (k + 1))

```

**Notation** is defined in STypes/Notation.lean following the mathematical conventions established in Section 2.3 and extended in LN Syntax for Types. The only noteworthy divergence is that the exponentials  $!A$  and  $?A$  are written  $!!A$  and  $??A$  respectively, to avoid clashing with Lean's built-in  $!$  (boolean negation) and  $?$  (syntactic hole) operators.

The full inductive definition of Types and all lemmas from STypes/Lemmas.lean are listed in ??.

**4.1.2 Processes.** Processes are defined in Processes/Basic.lean as the inductive type Proc. Following the established locally nameless architecture, channel references are partitioned into a dedicated Channel inductive. Explicit free channels are instantiated as concrete natural numbers

(FPName). While the raw process syntax permits arbitrary reuse of these numbers, using Nat allows the formalisation to systematically compute strictly fresh names by incrementing the supremum of any finite set of existing names (mechanised in Processes/Fresh.lean). Conversely, locally bound channels are represented purely by their relative scoping depth via an anonymous De Bruijn index (BPNam) (formalised also as a natural number, where, as well see in reference to the typing section showing how process level binders are calculated, their usage is restricted / constructed to avoid capture and automatically track binder distance):

```

inductive Channel : Type where
  | free  (x : FPName)
  | bound (x : BPNam)

```

As discussed in Section 3.2, binding constructs in the process syntax no longer carry explicit payload variables. Instead, bound occurrences are captured purely by their relative position. Comparing the standard paper grammar with the intermediate locally nameless notation and the final constructors of the inductive Proc type illustrates this abstraction directly:

| Standard              | Locally Nameless           | Lean Constructor                   | Binds      |
|-----------------------|----------------------------|------------------------------------|------------|
| $x[y].P$              | $x[\bullet].P$             | tensor (x : Channel) (P : Proc)    | 1 channel  |
| $x(y).P$              | $x(\bullet).P$             | parr (x : Channel) (P : Proc)      | 1 channel  |
| $\nu xy.P$            | $\nu(\bullet, \bullet).P$  | cut (P : Proc)                     | 2 channels |
| $x(X).P$              | $x(\diamond).P$            | input (x : Channel) (P : Proc)     | 1 type     |
| $x[\text{DUP}](x').P$ | $x[\text{DUP}](\bullet).P$ | duplicate (x : Channel) (P : Proc) | 1 channel  |

Notice that cut takes only a continuation  $P$  with no explicit channel arguments, as both dual endpoints are simultaneously bound as indices 0 and 1 within the scope of  $P$ . Similarly, for constructs like tensor, parr, and input, the subject channel  $x$  over which the interaction occurs is kept explicit, while the transmitted payload is anonymized into the continuation.

**Opening and closing** are defined in Processes/OpenClose.lean following the Locally Nameless principles established in Section 3.3.

Because the architecture relies on spatial tracking, the depth counter  $k$  increments according to the number of names bound by each constructor. For standard communication prefixes like tensor, parr, and the server duplicate operation, exactly one channel is bound, incrementing  $k$  by one. The restriction constructor (cut) simultaneously binds two dual endpoints, requiring  $k$  to increment by two. Notably, since processes can also bind types via input, where even though it is a binder, it does not bind a name, so we do not increment  $k$  in this case:

```

def Proc.open (P : Proc) (u : Channel) (k : Nat) : Proc :=
  match P with
  | .tensor x P      => .tensor (x.open u k) (P.open u (k + 1))
  | .parr   x P      => .parr   (x.open u k) (P.open u (k + 1))
  | .cut      P      => .cut      (P.open u (k + 2))
  | .input  x P      => .input  (x.open u k) (P.open u k)
  | ...

```

At the leaves of the syntax tree, this recursive traversal delegates to the base case defined by Channel.open. For bound channels, the function evaluates the De Bruijn index  $i$  against the target depth  $k$ , replacing it with the instantiated channel  $v$  upon an exact match. Explicit free channels remain entirely unaffected:

```

981 def Channel.open (u v : Channel) (k : Nat) : Channel :=
982   match u with
983   | bound i => if i == k then v else bound i
984   | c => c

```

For restriction, which requires instantiating two endpoints simultaneously, the `HasOpenTwo` typeclass from `Base.lean` is instantiated by opening sequentially: first replacing index  $k + 1$  with the second endpoint  $v$ , then replacing index  $k$  with the first endpoint  $u$ :

```

989 instance : HasOpenTwo Proc Channel Channel Nat where
990   open_ P u v k := (Proc.open (Proc.open P v (k + 1)) u k)

```

To complete the bidirectional infrastructure, `Proc.close` and `Channel.close` define the exact inverse operation. Used primarily when abstracting explicit free names into new scopes, `Proc.close` traverses the process tree using a structural recursion identical to `Proc.open`, applying the exact same depth-incrementing rules for single and double binders. When the recursion reaches a leaf node, it delegates to `Channel.close`. If the leaf is a free channel (`FPName`), the function compares it against the target name  $x$ ; upon a match, the explicit name is abstracted into a bound De Bruijn index at the accumulated depth  $k$ :

```

998 def Channel.close (u : Channel) (x : FPName) (k : Nat) : Channel :=
999   match u with
1000   | .free name => if x == name then .bound k else .free name
1001   | .bound i   => .bound i

```

**Local closure** for processes is defined in `Processes/LocalClosure.lean` as a two-parameter predicate `Proc.lc (k n : Nat) : Proc → Prop`. Because the calculus features a two-sorted syntax, the predicate must maintain independent depth counters:  $k$  tracks the depth of channel binders, while  $n$  tracks the depth of type binders. This ensures that when a process embeds a type, such as the payload  $A$  in an output operation, the embedded type can be correctly evaluated against the current type depth  $n$  during its own local closure check.

The structural recursion evaluates each constructor by incrementing the appropriate counter based on the number and sort of binders introduced. For instance, a tensor prefix binds a single continuation channel, incrementing the channel depth to  $k + 1$ . The restriction operation (`cut`) binds two dual endpoints simultaneously, incrementing the depth to  $k + 2$ . Conversely, an input prefix binds a polymorphic type variable rather than a channel. This leaves the channel depth  $k$  unchanged while incrementing the type depth to  $n + 1$ :

```

1015 def Proc.lc : Nat -> Nat -> Proc -> Prop
1016 | k, n, .tensor x P   => x.lc k ∧ P.lc (k + 1) n
1017 | k, n, .cut P        => P.lc (k + 2) n
1018 | k, n, .input x P    => x.lc k ∧ P.lc k (n + 1)
1019 | ...

```

A process is considered globally *closed* when it contains no dangling indices of either sort, meaning `Proc.lc 0 0` holds. This is registered as the abbreviation `Proc.lc_0`.

**Free names** are computed via `Processes/Names.lean` using the function `Proc.f : Proc → Finset FPName`, which traverses a process and collects all explicit channel identifiers into a finite set. Because locally nameless binders are represented by anonymous indices, they contribute nothing to this set. Consequently, constructors like `cut`, which operate entirely on bound endpoints and take no explicit subject channels, simply pass through the free names of their continuation:

```

1028 def Proc.f : Proc -> Finset FPName

```

```

1030 | .cut P      => P.f
1031 | .par P Q   => P.f ∪ Q.f
1032 | .link x y => x.f ∪ y.f
1033 | ...

```

**Freshness** is established in `Processes/Fresh.lean` to support the generation of globally unique names. Given a finite set  $L$  of forbidden names, a strictly fresh name is deterministically computed as  $\max(L) + 1$ . The formalisation provides two critical lemmas, `exists_one_fresh` and `exists_two_fresh`, which guarantee the generation of one or two distinct names residing entirely outside of  $L$ . These lemmas are essential for providing witnesses to satisfy the co-finite quantification premises used throughout the typing and transition rules, particularly for restriction, which requires introducing two fresh, non-colliding endpoints simultaneously.

**Name substitution** is defined in `Processes/SubstNames.lean` as `Proc.substNames(R T : FPNName) : Proc` replacing all free occurrences of name  $T$  with name  $R$ . Since bound channels are represented as indices, substitution never touches them and is therefore capture-avoiding by construction. Thus, no freshness conditions are needed.

**Type substitution** is defined in `Processes/SubstTypes.lean` as `Proc.substTypes (A : Types) (k : Nat)` substituting type  $A$  for the bound type index  $k$  throughout a process. The only constructor that carries a type directly is output, where `Types.subst` is applied to the embedded type. For input, which binds a type variable, the substitution descends with  $k + 1$  and the substituted type  $A$  is shifted by one to maintain correct index distances, mirroring the pattern established for `Types.subst` in ??:

```

1051 | .input x P => .input x (P.substTypes (A.shift 0 1) (k + 1))

```

The full inductive definition of `Proc` and all lemmas are listed in ??.

#### 4.1.3 Environments and Hyperenvironments.

#### 4.1.4 Typing Judgements.

### 4.2 Operational Semantics

The `Semantics` namespace covers the operational semantics of  $\pi$ MLL. Transition labels are defined in `Label.lean`. The LTS over typing derivations is defined in `JudgementStep/`, and the induced LTSs over processes and typing environments are derived in `ProcStep/` and `EnvStep/` respectively, realising the erasure homomorphisms described in [Section 2.5](#).

### 4.3 Session Fidelity

The `SessionFidelity` namespace establishes the results for  $\pi$ MLL. It contains the simulation results showing `EnvStep` and `ProcStep` can each mimic `JudgementStep`, subject reduction, which states that if  $P$  is well-typed and takes a step to  $P'$ , then  $P'$  is well-typed under a stepped environment, and typability, which lifts subject reduction to the level of typing derivations, showing that a well-typed process and its typing environment step together coherently.



---

**ITEMS    TODO**

|         |                                                                                                            |
|---------|------------------------------------------------------------------------------------------------------------|
| Writing | Turn Note paragraphs into remark environments                                                              |
| Writing | Remove highlighting from process names not being part of a process in Background and BindingRepresentation |
| Writing | Fix content spread between the Binding Representation and Formalisation sections                           |
| Writing | Hold off on explaining rule side conditions too deeply in background; defer to formalisation chapter       |
| Writing | In the metatheory/proof strategies section, discuss the structural vs. step induction dilemma for erasure  |
| LaTeX   | Fix <a href="#">Proposition 2.9</a> reference in the section on <code>Typing_f_eq_names</code>             |
| LaTeX   | Make appendix of the full inductives from Lean                                                             |

---

## 5 Future Work (WIP)

*Operational semantics.* While the current development provides a complete formalization of syntax, typing, and substitution for all constructs, the operational semantics and associated metatheory have been established for the multiplicative fragment not including the restriction rule. A natural direction for future work is to extend the operational semantics by completely incorporating the res-rule into the proofs of session fidelity. This would likely require additional reasoning about its behavior when interacting under an application of itself (basically an inversion lemma), and the any permutation its associated hyperenvironment might have, which could require the knowledge that a non-active component in a hyperenvironment is preserved across steps.

Extending the formalization to the full  $\pi$ LL system, would be the ultimate end goal, and would require the addition of a constructors for each of the rules in the non-multiplicative fragment of  $\pi$ LL to the already existing `TypingStep` inductive. An adaptation of all the lemmas pertaining to the properties of the `TypingStep` relation would need to be updated to account for these new constructors, and similar reasoning as for the res-rule is expected to be needed for each of the new rules to complete their respective cases in the proof for session fidelity.

Given that the foundational infrastructure for binding and substitution is already in place, this extension is expected to integrate smoothly albeit with the existing development maybe a little tediously when it comes to proving the various multi-thousand line permutations lemmas, which will likely be required. Completing these step would then result in a fully mechanized account of  $\pi$ LL, providing a tool for rigid validity checking of the already sound and existing theory

*Structural Congruence.* Another matter of interest would be addressing the mechanization of structural congruence. In the current implementation, the strict binary tree structure of the inductive types causes certain well-typed, theoretically reducible processes to become artificially stuck. This primarily manifests through the commutativity and associativity of parallel composition ( $P \mid_p Q \equiv Q \mid_p P$  and  $P \mid_p (Q \mid_p R) \equiv (P \mid_p Q) \mid_p R$ ), where the synchronization rule strictly expects interacting processes to be immediate siblings in a specific left-to-right syntactic order

Furthermore, mechanizing other standard congruences, such as Scope Extrusion, introduces unique challenges specific to the chosen syntax representation. By adopting a Locally Nameless (LN) approach, the traditional  $\alpha$ -equivalence and fresh-name side conditions of Scope Extrusion are elegantly bypassed. However, they are replaced by the complexities of De Bruijn index arithmetic. Pulling a parallel process  $Q$  under a restriction binder ( $\nu xy P \mid Q \equiv \nu xy P \mid (\text{shift } 0 \ 2 \ Q)$ ) requires explicitly shifting the bound indices within  $Q$  to account for the increased binder depth. Because the typing hyperenvironments are mechanized using a List-based representation, associativity and commutativity of the contexts must be managed through explicit structural permutation rules.

Formally proving that these De Bruijn index manipulations distribute and commute safely across explicit list permutations and structural merges adds substantial proof overhead.

While replacing binary parallel composition with unordered multisets of processes is a common theoretical remedy to these rigidities, implementing this in a dependent type theory like Lean introduces severe dependent elimination issues. Because multisets are formalized as quotients of lists, indexing inductive typing judgments over them destroys syntax-based matching and inversion, drastically complicating proof search. Consequently, future iterations would be best served by maintaining the rigid binary abstract syntax tree and list-based environments. While the existing permutation rules allow for necessary flexibility in the typing context (background), adding a process congruence rule would make the transition system significantly heavier by collapsing the 1-to-1 mapping between syntax and behavior. Instead, extending the semantics with explicit symmetric structural rules (e.g., syn-symm) allows transitions modulo congruence while maintaining the syntactic tractability and index stability essential for manageable proofs in Lean.

## References

- Brian Aydemir, Arthur Charguéraud, Benjamin C. Pierce, Randy Pollack, and Stephanie Weirich. 2008. Engineering Formal Metatheory. In *Proceedings of the 35th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages* (San Francisco, California, USA) (POPL '08). ACM, New York, NY, USA, 3–15. doi:10.1145/1328438.1328443
- Brian E. Aydemir, Aaron Bohannon, Matthew Fairbairn, J. Nathan Foster, Benjamin C. Pierce, Peter Sewell, Dimitrios Vytiniotis, Geoffrey Washburn, Stephanie Weirich, and Steve Zdancewic. 2005. Mechanized Metatheory for the Masses: The PoplMark Challenge. In *Theorem Proving in Higher Order Logics*, Joe Hurd and Tom Melham (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 50–65.
- Luís Caires and Frank Pfenning. 2010. Session Types as Intuitionistic Linear Propositions. In *CONCUR*. 222–236.
- Marco Carbone, Sam Lindley, Fabrizio Montesi, Carsten Schürmann, and Philip Wadler. 2016. Coherence Generalises Duality: A Logical Explanation of Multiparty Session Types. In *27th International Conference on Concurrency Theory, CONCUR 2016, August 23–26, 2016, Québec City, Canada (LIPIcs, Vol. 59)*, Josée Desharnais and Radha Jagadeesan (Eds.). Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 33:1–33:15. doi:10.4230/LIPIcs.CONCUR.2016.33
- Arthur Charguéraud. 2012. The Locally Nameless Representation. *Journal of Automated Reasoning* 49, 3 (2012), 363–408. doi:10.1007/s10817-011-9225-2
- N. G. de Bruijn. 1972. Lambda calculus notation with nameless dummies, a tool for automatic formula manipulation, with application to the Church-Rosser theorem. *Indagationes Mathematicae (Proceedings)* 75, 5 (1972), 381–392. doi:10.1016/1385-7258(72)90034-0
- Roberto Di Cosmo and Dale Miller. 2023. Linear Logic. In *The Stanford Encyclopedia of Philosophy* (fall 2023 ed.), Edward N. Zalta and Uri Nodelman (Eds.). Metaphysics Research Lab, Stanford University. <https://plato.stanford.edu/archives/fall2023/entries/logic-linear/>
- Jean-Yves Girard. 1987. Linear Logic. *Theor. Comput. Sci.* 50 (1987), 1–102. doi:10.1016/0304-3975(87)90045-4
- Robin Milner, Joachim Parrow, and David Walker. 1992. A Calculus of Mobile Processes, I. *Inf. Comput.* 100, 1 (1992), 1–40.
- Fabrizio Montesi and Marco Peressotti. 2021. Linear Logic, the  $\pi$ -calculus, and their Metatheory: A Recipe for Proofs as Processes. *CoRR* abs/2106.11818 (2021). arXiv:2106.11818 doi:10.48550/arXiv.2106.11818
- Philip Wadler. 2014. Propositions as sessions. *Journal of Functional Programming* 24, 2–3 (2014), 384–418. Also: ICFP, pages 273–286, 2012.